

Vidya - Landing Page Build Brief

Cold-start handoff. Written 2026-08-19 from the live codebase at `feat/v2-foundation` @ `e927857` . If you are a fresh session picking this up with no memory: **read §1 before writing a single line of copy**, then §2, then the rest top-to-bottom.

Deliverable: a high-converting marketing landing page for Vidya, deployed to Vercel, plus an interactive README for a new public marketing repo (spec'ed separately in `INTERACTIVE-README-SPEC.md`). Brand tokens, voice and the full copy deck are in `BRAND-AND-COPY-DECK.md` .

Owner decisions already made (do not re-litigate):

Decision	Choice
Primary CTA	Live demo + waitlist. "Try it now" (no signup) is primary; email capture is secondary.
Coverage honesty	Radical transparency. Ship an interactive coverage grid showing which chapters are deep and which are thin.
Public repo scope	Marketing/landing repo only. The app source stays private.

1. Hard constraints - read first

1.1 The landing page markets the `publicStudy` build. Nothing else.

Vidya's codebase builds three variants (`docs/v2/BUILD-VARIANTS.md`):

Variant	Private messaging	Public?
<code>publicStudy</code>	Never	Yes - this is what you are marketing
<code>trusted</code>	Yes, disclosed	No
<code>development</code>	Yes	No

The `trusted` / `development` surface must not appear anywhere in the landing page, the README, the repo, commit messages, alt text, meta tags, OG images, source comments, or `robots.txt` . No mention, no hint, no "and more", no knowing wink. This is not a stylistic preference:

- It is the project's own rule. `capabilities.ts` defines `PRIVATE_CAPABILITIES` and the existing onboarding feature is already asserted - in domain tests `*and*` against the rendered screen - to leak the private side in "no title, detail, id or destination". The landing page is held to the same standard.
- `publicStudy` genuinely is a study app. Marketing it as one is not a half-truth; it is the whole truth about that artifact.

Practical rule for the builder: work only from the facts in §2 of this doc. Do not go spelunking in `BLUEPRINT.md` , `OVERVIEW.md` , `PROJECT.md` or `KNOWLEDGE.md` for copy - those describe the full private

product and will contaminate the page. Everything you need to write truthful marketing is here.

1.2 Do not claim what is not true

The project's entire engineering culture is built on refusing unproven claims, and the landing page is not exempt. As of 2026-08-19 **all of these are false**:

Do NOT say	Reality
"Download on Google Play" / "Get the app"	No Play listing exists. No signed release. Stage 7 of the release gate is skipped for want of a keystore.
"Trusted by N students" / any usage number	Zero users. No analytics. No installs.
"Battle-tested on Android"	The APK builds locally; it has never been installed or launched on a device or emulator. adb has never been involved.
"Rated 4.8" / testimonials / logos	None exist. Do not invent, and do not use placeholder faces that imply real users.
"AI tutor built in"	The AI copilot is source-ready but not wired to a live provider - no key, no upstream, no deploy. Market it as *coming*, or omit it.
"Complete CBSE coverage"	50 of 62 chapters are thin. See §1.3.

Anything you are unsure about: leave it out. An empty section converts better than a claim a visitor can falsify in one tap.

1.3 The coverage problem, and the decision about it

3,903 questions across 62 chapters is arithmetically true and **materially misleading**: 12 chapters carry 300+ questions each, and the other 50 carry between 1 and 6. A student who lands, signs up, opens "Cash Flow Statement" and finds 3 questions has been mis-sold.

The owner has chosen radical transparency. Build the coverage grid (§5.2). Lead with depth where it exists, name the gap plainly, and frame it as a public roadmap. Suggested headline framing: **"12 chapters are exam-ready today. Here's exactly which."** This is a differentiator, not an apology - no competitor in this space publishes its own gaps.

2. Verified product facts

Every number below was generated from the live dataset on 2026-08-19, not from any doc. **Regenerate before launch** - see §7.1; do not hand-type these.

2.1 Headline numbers

Fact	Value
Questions in the bank	3,903
Chapters	62
Chapters at full depth (≥ 100 Easy / ≥ 100 Medium / ≥ 100 Hard)	12
Questions with a written explanation	3,895 (99.8%)
Questions with step-by-step worked solutions	3,408
Full-length mock papers	26
Duplicate questions	0
Data-integrity defects	0
Commerce calculators	10

2.2 By subject

Subject	Questions
Maths	1,590
Accountancy	1,267
Economics	679
Business Studies	367

2.3 By exam (a question can carry several tags)

Exam	Eligible questions	Mock papers
CUET	3,222	10
CA Foundation	2,036	8
IPMAT	1,851	8
CBSE Board	1,357	-
CMA Foundation	1,009	-
SET	727	-
NPAT	640	-
CS Foundation	435	-

2.4 The 12 exam-ready chapters

Accounting Ratios · Applied Maths (Commercial Arithmetic) · Company Accounts - Share Capital · Goodwill
 · Matrices & Determinants · Measures of Central Tendency · National Income & Aggregates · Partnership -

2.5 Features that genuinely exist and work

Verified in the running app, not aspirational:

- **Offline-first.** The whole bank ships in the app; no network needed to study. This is real and is the strongest honest differentiator.
- **Spaced repetition (SM-2)** over a due queue, seeded only by questions the student has actually attempted.
- **Practice** with instant feedback, written explanations and step-by-step working.
- **Mistake notebook** - wrong answers captured automatically.
- **Timed mock exams** with real negative marking, question palette, mark-for-review and single-submission guard.
- **Study planner** with per-day targets.
- **Progress analytics** - accuracy, streak, per-subject breakdown.
- **10 commerce calculators** (GST, interest, EMI, depreciation, ratios...).
- **Formula sheets & notes**, bundled and offline.
- **Getting-started checklist** that derives from the student's real stored work - it un-ticks if the work is cleared, and never fakes progress.

2.6 The quality story (your best trust asset)

This is unusually strong material and almost no ed-tech competitor can say it. Use it - it is all verifiable:

- Every numeric answer is ****re-derived** from the question text by an independent solver** that never reads the stored answer. If the content and the solver disagree, the build fails.
- Distractors encode **real misconceptions**, not noise - e.g. in Goodwill, the wrong options are "ignored the adjustment", "reversed its sign", and "forgot to multiply by years' purchase".
- Correct answers are distributed ~25% across each of the four positions, so there is no guessable pattern.
- **0 duplicates, 0 integrity defects**, enforced by an automated audit on every build.
- Nothing is labelled as a past paper unless its provenance is verified - fabricated "PYQ" claims are structurally impossible.

Suggested framing: **"Every number in this app is checked by a second program that solves the question independently. If they disagree, we don't ship."**

3. Positioning

Audience. Indian CBSE Class 11–12 Commerce students (and their parents), preparing simultaneously for Boards and one or more of CUET / CA Foundation / IPMAT. Price-sensitive, phone-first, frequently on poor or metered connections.

The insight. Commerce students are underserved. The ed-tech market is saturated with JEE/NEET products; Accountancy and Business Studies get scraps. Vidya's deepest content is exactly where the gap is.

Positioning statement.

For CBSE Commerce students preparing for Boards and entrance exams at the same time, Vidya is an offline-first practice app whose answers are machine-verified

- so you can trust every solution, on any connection.

Three pillars (use these as the page's spine):

1. **Works with no internet.** Not "offline mode" - the entire bank is on the device.
1. **Answers you can trust.** Independently verified; we publish the method.
2. **Built for Commerce, not adapted to it.** Accountancy and Maths are the deepest chapters.

Tone. Calm, precise, student-respecting. No hype, no exam-fear marketing, no "crack the exam" clichés, no countdown-timer pressure tactics. The honesty *is* the brand - the voice should sound like the person who wrote the audit.

4. Information architecture

Single scrolling page. Order is deliberate: proof before ask.

#	Section	Job
1	Hero + live question	Prove value in 10 seconds, before any ask (§5.1)
2	Three pillars	Offline · verified · Commerce-first
3	Coverage grid	Radical transparency; the trust moment (§5.2)
4	Exam picker	Personalise stats to CUET / CA / IPMAT (§5.3)
5	How verification works	The quality story from §2.6 - the moat
6	Feature tour	Practice, revision, mocks, mistakes, planner, calculators
7	Offline proof	The page itself survives going offline (§5.4)
8	Roadmap / honest status	What's built, what's next, no Play listing yet
9	Waitlist	Email capture, single field
10	Footer	Links, the public repo, contact

Mobile-first: the audience is on phones. Design at 375px, then scale up. The hero's playable question must work with one thumb.

5. The four unique features

These are the reason this page is not a template. Build all four; #1 and #2 are non-negotiable.

5.1 "Answer one right now" - playable hero

A real question from the real bank, rendered in the hero, answerable with no signup. On answer, reveal the real explanation and the step-by-step working.

- Pull from the 12 complete chapters; rotate per visit.
- **The wrong-answer state is the wow moment.** Because distractors encode

specific misconceptions, a wrong tap can say *"That's what you get if you forget to multiply by years' purchase."* Almost nothing else on the market can do this. Make the wrong state feel helpful, never punitive.

- Show the four-step working expanded after answering.
- Then, and only then, offer "Try 20 more →" (demo) or the waitlist.
- Reference pattern: [monday.com's "Try it out! See how monday.com works"](#)

and [The Leap's inline interactive demo](#).

5.2 The Honest Coverage Grid

A 62-cell grid, one cell per chapter, coloured by depth (deep / partial / thin), grouped by subject. Hover or tap a cell for exact counts. Generated at build time from the real audit (§7.1) - never hand-maintained.

- **Headline: "12 chapters are exam-ready. Here's exactly which."**

- Include a plain-language legend and a one-line note that thin chapters are being authored, with the current total.
- This is the page's trust centrepiece. Do not soften it, and do not let it drift out of sync with the dataset.

5.3 Exam picker

A segmented control - CUET · CA Foundation · IPMAT · Boards. Selecting one re-renders the eligible-question count, the mock-paper count and the relevant subjects, from §2.3. Instant personalisation, no form, no signup.

5.4 Offline proof

The landing page ships a service worker and works fully offline after first load. A small section invites the visitor to switch off their connection and keep scrolling; a live indicator confirms the page is offline and still working.

- This *demonstrates* the core product claim instead of asserting it.
- Must degrade gracefully: if the service worker fails to register, hide the invitation rather than making a claim the page can't keep.

6. Tech stack and repo

Stack. Next.js (App Router) + TypeScript + Tailwind on Vercel. Framer Motion for motion. No CMS - content is generated (§7.1). No analytics that sets cookies without consent; prefer a cookieless, privacy-respecting option or none at all (the brand is privacy-adjacent; a tracker wall would be self-defeating).

New public repo - marketing only. Suggested `vidya-landing`.

```
vidya-landing/  
  README.md           <- the interactive README (separate spec)  
  app/  
    page.tsx          <- the landing page  
    api/og/route.tsx   <- OG image  
    api/readme/stats/route.ts <- SVG endpoint for the README  
    api/readme/question/route.ts <- daily question SVG for the README  
  components/  
  content/  
    stats.generated.json <- generated; never hand-edited  
    questions.sample.json <- generated; a curated public subset  
  scripts/  
    generate-content.ts  
  public/
```

Never commit to the public repo: anything from `apps/mobile/src` beyond the generated public subset, any relay or crypto material, any reference to the private surface, `.env` values, keystores, or the full question bank (see §7.2).

7. Content pipeline

7.1 Stats are generated, never typed

This mirrors the project's core anti-drift discipline (every drift it has ever found came from a fact expressed twice). A hand-typed "3,903" on a landing page is that same bug.

Add `scripts/generate-content.ts` to the landing repo, run in `prebuild`, which emits `content/stats.generated.json` containing the §2 numbers plus the 62-chapter coverage array. Source it from the private repo's dataset - either by running the audit there and committing the JSON output across, or by vendoring a snapshot. **The page must import the generated file; a literal number in JSX is a bug.**

Add a check that fails the build if `stats.generated.json` is older than the dataset snapshot it was built from.

7.2 The public question subset

The hero and README question need real questions in a public repo. Do **not** publish the whole bank - it is the product.

Export a curated subset (suggest 40–60) drawn from the 12 complete chapters, with explanation and `step_solution` intact, into `content/questions.sample.json`. Note in the repo README that it is a sample, not the bank.

8. Definition of done

- ☐ Deployed to Vercel; custom domain configured.
- ☐ Lighthouse ≥ 95 performance / ≥ 95 accessibility on **mobile**.
- ☐ Fully usable at 375px width with one thumb; no horizontal overflow at 320px.
- ☐ Hero question is answerable, and the explanation + working render correctly.
- ☐ Coverage grid renders all 62 chapters from generated data.
- ☐ Exam picker updates all four stats correctly.
- ☐ Page works offline after first load; the offline invitation hides if the SW fails.
- ☐ Zero hand-typed statistics anywhere in the source.
- ☐ **Zero references to the private surface** anywhere in the repo, build

output, or metadata. Grep the built output for the terms in §1.1 before shipping and make that grep a CI step.

- ☐ Every factual claim traceable to §2 of this doc.
 - ☐ No fabricated testimonials, ratings, user counts, or partner logos.
 - ☐ Respects `prefers-reduced-motion`; all interactive elements keyboard-reachable.
 - ☐ OG image and meta tags set; social preview checked.
-

9. Open items for the owner

1. **Domain** - not chosen.

2. **Live demo target.** The hero's "Try it now" should point at a deployed

`publicStudy` **web** build. That build does not exist yet as a public deployment. Until it does, the CTA should scroll to the waitlist rather than

1. Note: `DEPLOY-WEB.md` describes deploying the *full* app with a relay - that is the wrong artifact here; a `publicStudy` web build needs no relay.

1. **Waitlist backend** - no provider chosen. Needs to store an email and nothing else; state the retention policy on the form.

1. **App name/trademark check** - "Vidya" is a common word; worth a search before buying a domain.